

Databases Project – Spring 2021

Team No: 65

Members: Zad ABI FADEL, Loïc HOUMARD, Jonas BLANC

Contents

| | |
|---|----|
| Contents..... | 1 |
| Deliverable 1..... | 2 |
| Assumptions | 2 |
| Entity Relationship Schema | 2 |
| Schema..... | 3 |
| Description..... | 3 |
| Relational Schema | 4 |
| ER schema to Relational schema | 5 |
| DDL..... | 5 |
| General Comments | 6 |
| Deliverable 2..... | 15 |
| Assumptions | 16 |
| Data Loading/Cleaning..... | 16 |
| Query Implementation | 18 |
| General Comments | 30 |
| Deliverable 3..... | 31 |
| Assumptions | 31 |
| Query Implementation | 31 |
| Query Performance Analysis – Indexing | 32 |
| General Comments | 61 |

Deliverable 1

Assumptions

On Identification:

Every party number should be unique within a collision. Every `party_id`, `victim_id`, `case_id` should be unique by its own within the corresponding .csv files.

On data:

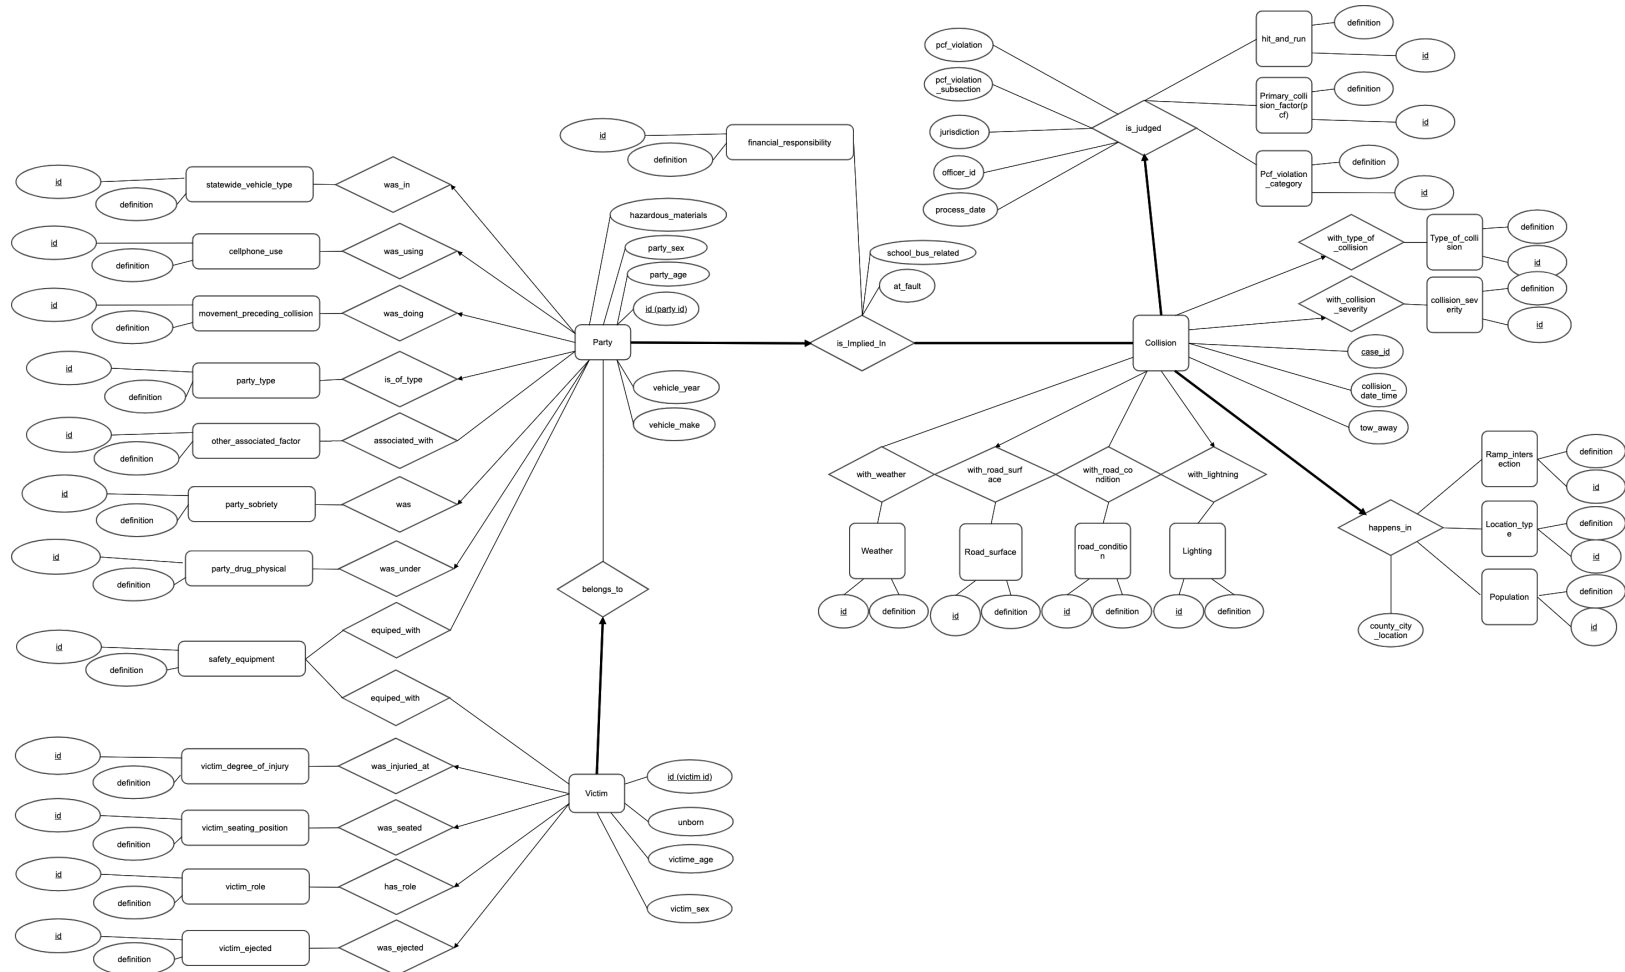
We assumed that in the .csv files every field would be represented by its key or that we would make it so during the data cleaning phase. We assumed that every description could fit in 150 char. We assumed based on data that `party_id`, `victim_id` and `case_id` can be typed as integer.

On integrity:

Every victim should be associated with an unique party. Every party should be implicated in a unique collision.

Entity Relationship Schema

Schema



Description

For the ER diagram, we first decided to divide the attributes into 3 main entities called Victim, Party and Collision, because it seemed to us that they were the main actors in the model.

Then, we saw that it didn't make much sense to have only these 3 entities, because some attributes wouldn't be logically attributed to them. For example, it wouldn't make sense that a collision has an attribute population, because they are not directly correlated. Therefore, we tried to group attributes that logically belonged to a common idea together (star schema). For the collisions, we saw that there were many attributes related to the location of the collision, the conditions under which the collision happened and the legal part related to the collision. For the parties, many attributes were related to the vehicle. Hence, we wanted to add these 4 entities to our diagram (but finally modified it slightly, see below).

Also, after we spoke with some assistants, we realised that it would be a good idea to create entities for attributes that are lists with some finite non-logically predefined values (A:..., B:...). The reasons are the following: it would be easier to enforce the data we store to be cleaned and in the same format (it avoids to have one time 'a' and one time 'A' referencing to the same value) and it would make it more modifiable and easier to change (if we realize that we would like to add/remove an option, we could simply add/remove one row in the table of the entity and add/invalidate these entries in the other table).

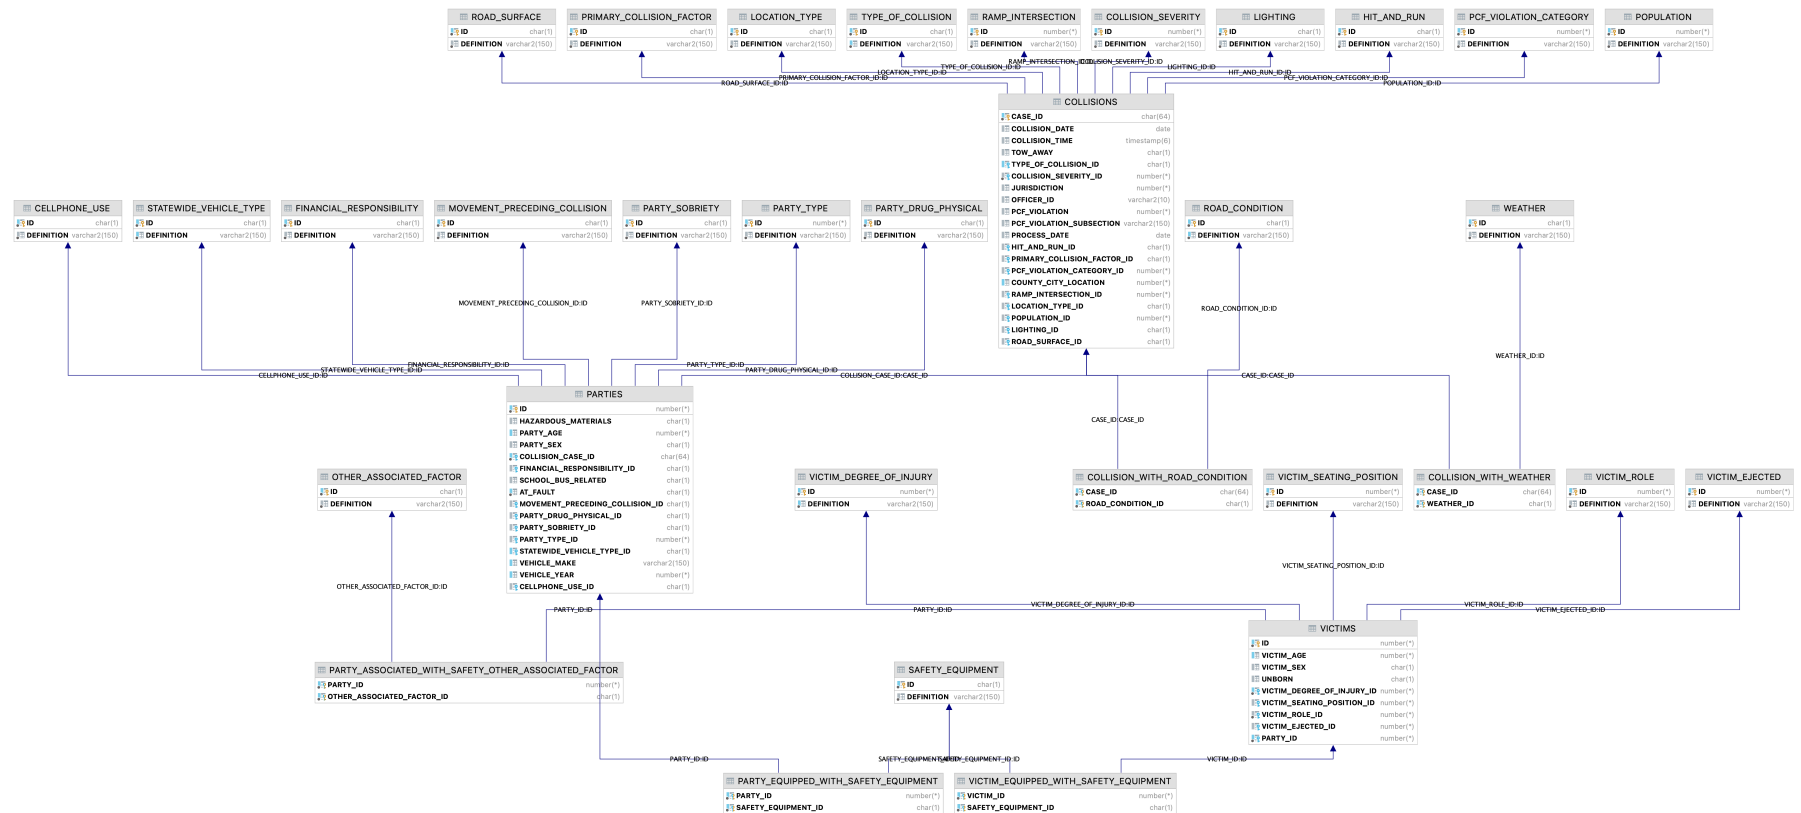
When there were many times the same attribute in the csv files (..._1 and ..._2), we also decided to create an entity. This has the advantage to be more modifiable, since we could decide to add a third (..._3) attribute or even more of them in the future if we would like to slightly change the model. For that, we simply allowed the relation to have many of these new entities.

Finally, when we wanted to merge all our previous ideas together to construct the diagram, we found that creating the 4 entities mentioned above was not really practical because we would have to create these entities which now have no (or not many) attributes (since their corresponding attributes were often lists which we now model with an entity and bind through a relation), which makes them almost useless and increases the complexity of the diagram. Therefore, we decided to create N-ary relations directly to group the collision and all the attributes related to a given theme. This seems easier to understand and will create the same result in the database (since every attribute will finally be stored in the Collision table after the merging due to the many-to-one relation) when we translate it from the ER model to the SQL DDL commands.

After the first milestone, we also decided to remove the condition table which we had kept, because we found it easier to implement in the data cleaning process and because our associated TA advised us to do so. Indeed, on our older schema, we had to create a custom key for condition and bind it through a relation which was more complicated and didn't bring much. The only utility of the condition table was to make the star schema easier to understand, but in practice it didn't bring much.

Relational Schema

ER schema to Relational schema



DDL

```
---Design implementations---
-- Boolean => char(1)
-- definition => varchar(150)
-- Table_name (First letter upper case then underscores)
-- One-to-Many (Store key in one)
-- No state is null, set key to null
-- In an entity: id is id of current entity, create new attribute
table_id for referenced id

---Collisions start---
CREATE TABLE Weather
(
    id          char(1), -- check if id is one of letter
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Road_surface
(
    id          char(1), -- check if id is one of letter
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Road_condition
(
    id          char(1), -- check if id is one of letter
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Lighting
(
```

```
    id          char(1), -- check if id is one of letter
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Type_of_collision
(
    id          char(1), --check char between a & h
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Collision_severity
(
    id          int CHECK (0 <= id and id <= 4),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Hit_and_run
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Primary_collision_factor
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Pcf_violation_category
(
```

```
    id          int CHECK ((0 <= id and id <= 24)),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Ramp_intersection
(
    id          int CHECK (1 <= id and id <= 8),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Location_type
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Population
(
    id          int CHECK (0 <= id and id <= 9),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Collisions
(
    case_id          char(64),
    collision_date    date,
    collision_time    timestamp(6),
    tow_away         char(1) CHECK (tow_away = 'T' or
tow_away = 'F'),
    type_of_collision_id char(1) references
Type_of_collision (id),
```



```

    collision_severity_id      int not null references
Collision_severity (id),
    -- Relations is_judged
    jurisdiction              int CHECK (0 <= jurisdiction and
jurisdiction <= 9999),
    officer_id                varchar(10),
    pcf_violation              int,
    pcf_violation_subsection   varchar(150),
    process_date               date,
    hit_and_run_id            char(1) references Hit_and_run
(id),
    primary_collision_factor_id char(1) references
Primary_collision_factor (id),
    pcf_violation_category_id  int references
Pcf_violation_category (id),
    -- Relations happens_in
    county_city_location       int,
    ramp_intersection_id        int references Ramp_intersection
(id),
    location_type_id           char(1) references Location_type
(id),
    population_id              int references Population (id),
    -- Relations happens_under
    lighting_id                char(1) references Lighting (id),
    road_surface_id            char(1) references Road_surface
(id),
    PRIMARY KEY (case_id)
);

CREATE TABLE Collision_with_weather
(
    case_id      char(64) references Collisions (case_id) on delete
cascade,
    weather_id char(1) references Weather (id) on delete cascade,
    PRIMARY KEY (case_id, weather_id)
);

```

```
CREATE TABLE Collision_with_road_condition
(
    case_id          char(64) references Collisions (case_id) on
delete cascade,
    road_condition_id char(1) references Road_condition (id) on
delete cascade,
    PRIMARY KEY (case_id, road_condition_id)
);

---Collisions end---

CREATE TABLE Safety_equipment
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

---Parties start---

-- Related entities with party: one to many
CREATE TABLE Movement_preceding_collision
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Party_drug_physical
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);
```

```
CREATE TABLE Party_sobriety
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Party_type
(
    id          int,
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Statewide_vehicle_type
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Cellphone_use
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

-- Relations with party: Many to many
CREATE TABLE Other_associated_factor
(
    id          char(1),
    definition varchar(150) not null,
    PRIMARY KEY (id)
```

```
);

CREATE TABLE Financial_responsibility
(
    id          char(1),
    definition  varchar(150) not null,
    PRIMARY KEY (id)
);

-- Parties
CREATE TABLE Parties
(
    id                      int,
    -- Attributes
    hazardous_materials    char(1),
    party_age              int,
    party_sex              char(1),
    -- relation to collision
    collision_case_id       char(64) not null references
Collisions (case_id),
    financial_responsibility_id char(1) references
Financial_responsibility (id),
    school_bus_related     char(1),
    at_fault              char(1) not null,
    -- referenced ids
    movement_preceding_collision_id char(1) references
Movement_preceding_collision (id),
    party_drug_physical_id  char(1) references
Party_drug_physical (id),
    party_sobriety_id       char(1) references
Party_sobriety (id),
    party_type_id          int references Party_type (id),
    statewide_vehicle_type_id char(1) references
Statewide_vehicle_type (id),
    vehicle_make           varchar(150),
    vehicle_year           int,
```

```
    cellphone_use_id          char(1) default 'D' references
Cellphone_use (id), --default 'D' makes it faster
    -- key
    PRIMARY KEY (id)
);

CREATE TABLE Party_equipped_with_safety_equipment
(
    party_id          int      not null references Parties (id) on
delete cascade,
    safety_equipment_id char(1) not null references
Safety_equipment (id) on delete cascade,
    PRIMARY KEY (party_id, safety_equipment_id)
);

CREATE TABLE Party_associated_with_safety_other_associated_factor
(
    party_id          int      not null references Parties
(id) on delete cascade,
    other_associated_factor_id char(1) not null references
Other_associated_factor (id) on delete cascade,
    PRIMARY KEY (party_id, other_associated_factor_id)
);
---Parties end---

---Victims start---
CREATE TABLE Victim_degree_of_injury
(
    id          int CHECK (0 <= id and id <= 7), -- can we make sure
id and def are consistent
    definition varchar(150) not null,
    PRIMARY KEY (id)
);

CREATE TABLE Victim_seating_position
(
```



```
    PRIMARY KEY (id)
);

CREATE TABLE Victim_equipped_with_safety_equipment
(
    victim_id          int      not null references Victims (id) on
delete cascade,
    safety_equipment_id char(1) not null references
Safety_equipment (id) on delete cascade,
    PRIMARY KEY (victim_id, safety_equipment_id)
);
---Victims end---
```

General Comments

In general, we found it pretty hard to create the ER diagram at first because there were a lot of attributes to proceed and understand and also because we didn't have much experience with this kind of work. But after having spent some time, we think that our implementation is now logical and should allow us to retrieve the information without having too many problems.

The allocation between the members was good, since we almost always worked together as a team. We first all took part in the elaboration of the ER diagram by concentrating us each on a CSV file and then talking with each other to see which attributes could belong together. We then all wrote some of the SQL DDL commands to create the tables and wrote the report together.

Deliverable 2

Assumptions

Data Loading/Cleaning

We decided to clean the data in jupyter notebooks using pandas. We processed the data CSV by CSV then transferred the data using pickles for example to infer `party_id` from `case_id` and `party_number`. We used translation tables (python dictionary) to translate from description to id where it was needed, since we decided to create small entities for each for attributes that are lists with some finite non-logically predefined values. We generated the tables for such small entities by copying the data from the handout pdf file. For the relations with entities representing multiple attributes with the same mapping (with `_1` and `_2`) we concatenate all the non null rows and drop the duplicates since they don't add any information.

Collisions.csv:

No major assumptions were needed to clean the collisions data. We chose to use `timestamp` as a type for all the date and time attributes. We first wanted to use a specific type for date only and one for time only, but we didn't see any such data type available with Oracle DB, therefore we chose `timestamp` which is not ideal for our use case. For the `collision_date` the time is automatically set to 00:00. For the `collision_time` field we chose to set a fixed default date (2000-01-01) . We couldn't merge both date and time in a single field because when one of them is missing, setting it to a default value would compromise the integrity of the data.

`officer_id:`

We decided to change the officer id ",66" to None because we had problems inserting it in the database due to the ','. We could have changed it to "66" (which is a valid value in the dataset), but since we were not sure that it was a typo, we found this assumption too strong and therefore we preferred to remove it.

Parties.csv

The data from parties had more dirty values. Here are the choices we did:

`cellphone_use:`

We realised that the values that are stored in the `cellphone_use` column {'1', '2', '3', 'B', 'C', 'D', nan} are different to the ones on the handout {'B', 'C', 'D', nan}. The values that are in the data but not in the handout {'1', '2', '3'} appear 2'636'894 times. We decided not to drop these values because they are a big chunk of the data (56%).

We needed to find a plausible mapping between the numbers and the letters. We opted to do it by doing a frequency analysis.

| | |
|--------------------------|--------------------------|
| 1 : 24787 in % : 0.009 | B : 38932 in % : 0.018 |
| 2 : 39114 in % : 0.015 | C : 795475 in % : 0.377 |
| 3 : 2572993 in % : 0.976 | D : 1274423 in % : 0.604 |

As you can see, it is clear that 1 and B are those that appear the least, and 3 and D are those that appear most frequently.

Therefore, we concluded that the correct mapping is: 1 -> B, 2 -> C, 3 -> D

As we imported the data in the database we chose to replace the None value by "D" since D already means "No Cell Phone/Unknown" which is equal to "no value".

`vehicle_make:`

Since `vehicle_make` is an open field there are a lot of errors and inconsistency. We corrected the most obvious typos (see below) and made some brands consistent. We chose not to modify this field too much since we are not experts in `vehicle_make` and that's error prone to modify it manually. For example we decided not to remove values with "OTHER - ..." since they add information compared to a "None". Here are the typos and inconsistencies we corrected and:

| | |
|--------------------------------|----------------------------|
| "AMERICAN MOTORS" | => "AMERICAN MOTORS (AMC)" |
| "DODG" | => "DODGE" |
| "HOND" | => "HONDA" |
| "MERCEDES BENZ" | => "MERCEDES-BENZ" |
| "MAZD" | => "MAZDA" |
| "TOYTA" | => "TOYOTA" |
| "MISCELLANEOUS" , "NOT STATED" | => None |

`party_drug_physical:`

We noticed 585'062 rows of `party_drug_physical` with value "G" which is not a valid key. We decided to replace it by `None` since we had no way to guess what the correct value was.

Victims:

`victim_age` and `pregnancy:`

In order to clean the data and make querying easier, we decided to create a new field: `unborn` which is a boolean telling if the victim was born or not. We set `unborn` from the convention saying that if the age is a 999 then the victim is the fetus of a pregnant woman. Then we replaced the age 999 by `None`. We chose to replace it by `None` and not 0 because we thought it would make more sense and that it would be weird if the mean of age of a 30 years old pregnant woman is 15 years.

Assumptions For Queries

For the queries, we assumed that we could use all available built-in functions for Oracle database systems. These functions are EXTRACT, COUNT, MEDIAN, FETCH, TO_CHAR, LOWER and DUAL.

Query Implementation

Description of logic:

This query should retrieve the number of collisions per year. Therefore, we first group by the year that we extract with the built-in function “EXTRACT(YEAR from ...)”. We then count the number of entries per year. We decided to order it by year, ascending to make it clearer.

SQL statement

```
SELECT EXTRACT(YEAR FROM C.COLLISION_DATE) AS YEAR, COUNT(*) AS  
NUMBER_COLLISIONS  
FROM COLLISIONS C  
GROUP BY EXTRACT(YEAR FROM C.COLLISION_DATE)  
ORDER BY EXTRACT(YEAR FROM C.COLLISION_DATE) ASC;
```

Query result (if the result is big, just a snippet)

| YEAR | NUMBER_COLLISIONS |
|------|-------------------|
| 2001 | 522562 |
| 2002 | 544741 |
| 2003 | 538954 |
| 2004 | 538295 |
| 2005 | 532725 |

| | |
|------|--------|
| 2006 | 498850 |
| 2007 | 501908 |
| 2017 | 7 |
| 2018 | 21 |

Query 2:

Description of logic:

This query should retrieve the most popular vehicle make and the number of vehicles for this make. We do this by first grouping by make and sorting it by the number of vehicles for each make. To retrieve the most popular make only, we use the “FETCH FIRST 1 ROW ONLY” built-in function (which is equivalent to limit in MySQL).

SQL statement

```
SELECT P.VEHICLE_MAKE, COUNT(*) AS NUMBER_VEHICLE
FROM PARTIES P
GROUP BY P.VEHICLE_MAKE
ORDER BY COUNT(*) DESC
FETCH FIRST 1 ROW ONLY;
```

Query result (if the result is big, just a snippet)

| VEHICLE_MAKE | NUMBER_VEHICLE |
|--------------|----------------|
| FORD | 1129701 |

Query 3:

Description of logic:

This query should retrieve the fraction of collisions which happen under dark lighting. For that, we first query the lightning that contains “dark” in their definition (note that we could directly use the ID since we know it, but we found clearer and more robust to query it using the definition if we would like to use the same query later on, when the table could be modified and more than one field could be about dark weather). We then bind it to the `lightning_id` stored in the collisions to count all the collisions with this weather type. We finally divide by the total number of collisions to have a fraction. We also decided to round the result to avoid having many useless digits.

SQL statement

```
SELECT
ROUND(A.NUMBER_COLLISIONS_UNDER_DARK/A.TOTAL_NUMBER_COLLISIONS, 3)
AS FRACTION_UNDER_DARK
FROM (
    SELECT
        (SELECT COUNT(*)
         FROM COLLISIONS C
         WHERE C.LIGHTNING_ID IN
            ( SELECT L.ID
              FROM LIGHTNING L
              WHERE LOWER(L.DEFINITION) LIKE '%dark%')) AS
NUMBER_COLLISIONS_UNDER_DARK,
        (SELECT COUNT(*) FROM COLLISIONS) AS
TOTAL_NUMBER_COLLISIONS
    FROM DUAL
) A;
```

Query result (if the result is big, just a snippet)

| FRACTION_DARK |
|---------------|
| 0.28 |

Query 4:

Description of logic:

This query should retrieve the number of collisions which happen under snowy weather. Just like before, we just query the ids in weather which contain “snow” in their definition and count all the entries of the relation which have this id.

SQL statement

```
SELECT COUNT(*) AS NUMBER_COLLISIONS_SNOWY_WEATHER
FROM COLLISION_WITH_WEATHER CWW
WHERE CWW.WEATHER_ID IN
    (
        SELECT W.ID
        FROM WEATHER W
        WHERE LOWER(W.DEFINITION) LIKE '%snow%');
```

Query result (if the result is big, just a snippet)

| NUMBER_COLLISIONS_SNOWY_WEATHER |
|---------------------------------|
| 8530 |

Query 5:

Description of logic:

This query should retrieve the number of collisions that happen every day of the week. For that, we first group by the day using the built-in function “TO_CHAR(date, ‘DAY’)” and count the number of entries. To retrieve the top 1 only, we first sort by the number of collisions and fetch the first row only.

SQL statement

```
SELECT TO_CHAR(C.COLLISION_DATE, 'DAY') AS WEEKDAY, COUNT(*) AS  
NUMBER_COLLISIONS  
FROM COLLISIONS C  
GROUP BY TO_CHAR(C.COLLISION_DATE, 'DAY')  
ORDER BY COUNT(*) DESC  
FETCH FIRST 1 ROW ONLY;
```

Query result (if the result is big, just a snippet)

| WEEKDAY | NUMBER_COLLISIONS |
|---------|-------------------|
| FRIDAY | 614853 |

Query 6:

Description of logic:

This query should retrieve all the types of weather and their corresponding number of collisions, sorted in descending order. For that, we simply join the tables weather and collisions with weather, then group by the definition and count all the entries.

SQL statement

```
SELECT W.DEFINITION, COUNT(*) AS NUMBER_COLLISIONS
FROM WEATHER W, COLLISION_WITH_WEATHER CWW
WHERE W.ID=CWW.WEATHER_ID
GROUP BY W.DEFINITION
ORDER BY COUNT(*) DESC;
```

Query result (if the result is big, just a snippet)

| DEFINITION | NUMBER_COLLISIONS |
|------------|-------------------|
| Clear | 2941042 |
| Cloudy | 548250 |
| Raining | 223752 |
| Fog | 21259 |
| Wind | 13952 |
| Snowing | 8530 |
| Other | 6960 |

Query 7:

Description of logic:

This query should retrieve all the `at_fault` collision parties with financial responsibility and loose material. For that, we had to check if the party is at fault in the table `party` and then check for the financial responsibility by using the ID and extracting the ones having “yes” in their description and finally check the loose material by using the `case_id` to retrieve the collision, then the `road_condition_id` from the relation table and finally take only the definition having “loose material” in it.

SQL statement

```
SELECT COUNT(*) AS NUMBER_AT_FAULT_WITH_FIN_REP_LOOSE_MAT
FROM PARTIES P, FINANCIAL_RESPONSIBILITY FR, COLLISIONS COL,
COLLISION_WITH_ROAD_CONDITION CWRC, ROAD_CONDITION RC
WHERE P.AT_FAULT = 'T'
AND P.FINANCIAL_RESPONSIBILITY_ID = FR.ID
AND LOWER(FR.DEFINITION) LIKE '%yes%'
AND P.COLLISION_CASE_ID = COL.CASE_ID
AND COL.CASE_ID = CWRC.CASE_ID
AND CWRC.ROAD_CONDITION_ID = RC.ID
AND LOWER(RC.DEFINITION) LIKE '%loose material%';
```

Query result (if the result is big, just a snippet)

| NUMBER_AT_FAULT_WITH_FIN_REP_LOOSE_MAT |
|--|
|--|

| |
|------|
| 4803 |
|------|

Query 8:

Description of logic:

This query should retrieve the median age and the most common victim seating position. Since these 2 pieces of information have not much to do with each other, we first wrote them individually and then used dual to write them together.

For the median age, we just used the built-in “MEDIAN” function.

For the most common victim seating position, we used the same trick as in query 2 which is to group by the seating position, sort by the number and keep the top row only.

SQL statement

```
SELECT
A.VICTIM_AGE_MEDIAN, A.MOST_COMMON_VICTIM_SEATING_POSITION

FROM
(
    SELECT
        (
            SELECT MEDIAN(V.VICTIM_AGE)
            FROM VICTIMS V) AS VICTIM_AGE_MEDIAN,
        (
            SELECT VSP.DEFINITION
            FROM VICTIM_SEATING_POSITION VSP
            WHERE VSP.ID IN
                (
                    SELECT V.VICTIM_SEATING_POSITION_ID
                    FROM VICTIMS V
                    GROUP BY V.VICTIM_SEATING_POSITION_ID
                    ORDER BY COUNT(*) DESC
                    FETCH FIRST 1 ROW ONLY)) AS
MOST_COMMON_VICTIM_SEATING_POSITION
        FROM DUAL
    ) A;
```

Query result (if the result is big, just a snippet)

| VICTIM_AGE_MEDIAN | MOST_COMMON_VICTIM_SEATING_POSITION |
|-------------------|-------------------------------------|
| 25 | Passengers |

Query 9:

Description of logic:

This query should retrieve the fraction of victims who were using a belt along all the participants. For that, we first count all victims which have a belt and divide by the total number of victims and participants using DUAL to be able to divide them. We also decided to round the result to make it more readable.

Remarks

We found this query not very logical since a party represents a group of people and that a party could be already counted in the victim table, but not necessarily since we have no way to be sure whether a party only has victims or not. At first, we had only counted the total number of victims (instead of victims + parties), but after seeing this post <https://moodle.epfl.ch/mod/forum/discuss.php?d=56137>, point3, we decided to use the query shown below.

SQL statement

```
SELECT ROUND(A.NUMBER_VICTIM_WITH_BELT / (A.TOTAL_VICTIM +
A.TOTAL_PARTIES), 3) AS FRACTION_WITH_BELT
FROM
(
    SELECT
        (SELECT COUNT(*)
         FROM VICTIMS V
         WHERE V.ID IN
            ( SELECT VEWSE.VICTIM_ID
              FROM VICTIM_EQIPPED_WITH_SAFETY_EQUIPMENT VEWSE
              WHERE VEWSE.SAFETY_EQUIPMENT_ID IN
                  ( SELECT SE.ID
                    FROM SAFETY_EQUIPMENT SE
                    WHERE LOWER(SE.DEFINITION) LIKE '%belt
use%' ))) AS NUMBER_VICTIM_WITH_BELT,
        (SELECT COUNT(*) FROM VICTIMS) AS TOTAL_VICTIM,
        (SELECT COUNT(*) FROM PARTIES) AS TOTAL_PARTIES
        FROM DUAL
    ) A;
```

Query result (if the result is big, just a snippet)

| FRACTION_WITH_BELT |
|--------------------|
|--------------------|

| |
|-------|
| 0.011 |
|-------|

Query 10:

Description of logic:

This query should retrieve the fraction of collisions that happen each hour of the day. For that, we simply group by the hour that we extract from the time using the `EXTRACT(HOUR, time)` built-in function, count the number of entries for each hour and divide by the total number of collisions.

Remark:

We decided to keep an entry when the hour was not specified with the fraction of accidents when the hour was unknown because we found it clearer this way.

We only showed the first 20 entries in the result as asked in the question.

SQL statement

```
SELECT EXTRACT(HOUR FROM C.COLLISION_TIME) AS HOUR,  
ROUND(COUNT(*) / (SELECT COUNT(*) FROM COLLISIONS), 3) AS  
FRACTION_COLLISIONS  
FROM COLLISIONS C  
GROUP BY EXTRACT(HOUR FROM C.COLLISION_TIME)  
ORDER BY EXTRACT(HOUR FROM C.COLLISION_TIME) ASC;
```

Query result (if the result is big, just a snippet)

| HOUR | FRACTION_COLLISIONS |
|------|---------------------|
| 0 | 0.019 |
| 1 | 0.018 |
| 2 | 0.018 |
| 3 | 0.012 |
| 4 | 0.01 |

| | |
|----|-------|
| 5 | 0.014 |
| 6 | 0.026 |
| 7 | 0.052 |
| 8 | 0.052 |
| 9 | 0.041 |
| 10 | 0.042 |
| 11 | 0.049 |
| 12 | 0.058 |
| 13 | 0.058 |
| 14 | 0.065 |
| 15 | 0.077 |
| 16 | 0.073 |
| 17 | 0.079 |
| 18 | 0.063 |
| 19 | 0.044 |

General Comments

We didn't have to change our previous work on the ER diagram in part 1 too much and we were able to write the queries quite easily. However, it took us a lot of time to clean the data and we had some problems when we tried to import the data in the database.

We decided to work all together on the different tasks, each team member spent an equal amount of time.

Deliverable 3

Assumptions

The assumptions/choices we made for our queries are the following:

- Round the decimal values returned by the queries to 3 decimal numbers to make them more readable.
- We decided to use the definitions of the small tables instead of directly using their ID because we found this way of querying the information easier to understand and cleaner. However, it comes with some cost since we must join the small tables everytime.
- For the 3rd query, we decided to discard the vehicle makes which were null. The reason is that null was the fourth most represented “vehicle make” and we found that this information was not really relevant since we were looking for real vehicle make (if a brand would like to make some statistics or know where they are on the list, they wouldn’t care about the null values which don’t give much useful information).
- For the 4th query, we understood “fraction of total incidents” as the number of incidents where no injury happened for a given seating position divided by the total number of victims seated at this particular position only (and not the total number of victims). It made more sense to us and a post on the forum seemed to agree with this assumption.
- For the 6th query, since many cities had the same population type (over 250’000) and we couldn’t know the exact population from the data, we just took the 3 first results that the database returned for this category.
- For the 6th query, we decided to keep the case_ids where some ages were unknown (null), but not considering these ages in the computation of the average. This means for example that if we have an accident with people of age (10,40, null), the average would be 25 since null would be ignored. We could have dropped these entries instead of accepting them and ignore the null values for computation only, but we found that it was a good approach to count the accidents where only partial values were given as well to limit the data we’re dropping. However, depending on why we would like to know this query, it could be useful to discard these entries.
- For the 8th query, since we didn’t use any vehicle id in our diagram, we decided to use the vehicle type, vehicle make and vehicle year as the id, because they represent all the available information we have about the vehicles.
- For the last query (10th), we first based our classification on the lighting information when they were clear enough (daylight for day, dark for night). For the dusk-dawn case, we based ourselves on the time and the month and when we had inconsistent data (dusk-dawn at 12:00 for example), we discarded it. When the lighting was null, we tried to infer the period based on the time and the date only when it was possible and discarded the data otherwise.

Query Implementation

Query 1:

Description of logic:

This query should retrieve the ratio of cases where the driver was at fault for different age groups. For this, we first did two subqueries: a first one which counts all the parties that fall in each age category where we simply discarded all the parties having null for age and grouped by the age category using a case on the age and a second one quite similar where we took only the parties being at fault for each age group following the same logic. We then returned the age category and divided each of the two results to have the ratio of parties at fault. We also decided to sort it in descending order to have a better vision of the results and clearly see which are the categories that were the most often at fault. As we can see in the results and as we might have expected, underage, young people and elder people (elder 2) tend to be more often at fault. Therefore, as an insurance company, it would make sense to make young and old people pay more than middle-aged adults for their insurance.

SQL statement

```
SELECT FAULT.age_range, ROUND(NUMBER_AT_FAULT / TOTAL_NUMBER, 3)
as RATIO_AT_FAULT
FROM (SELECT case
        when P.PARTY_AGE <= 18 then 'Underage'
        when P.PARTY_AGE between 19 and 21 then 'young 1'
        when P.PARTY_AGE between 22 and 24 then 'young 2'
        when P.PARTY_AGE between 24 and 60 then 'adult'
        when P.PARTY_AGE between 61 and 64 then 'elder 1'
        when P.PARTY_AGE >= 65 then 'elder 2' end as
age_range,
        COUNT(*) AS
NUMBER_AT_FAULT
FROM PARTIES P
WHERE P.AT_FAULT = 'T'
and P.PARTY_AGE IS NOT NULL
group by (case
        when P.PARTY_AGE <= 18 then 'Underage'
        when P.PARTY_AGE between 19 and 21 then 'young
1'
        when P.PARTY_AGE between 22 and 24 then 'young
2'
        when P.PARTY_AGE between 24 and 60 then 'adult'
        when P.PARTY_AGE between 61 and 64 then 'elder
1'
        when P.PARTY_AGE >= 65 then 'elder 2'
```



```

    END)) FAULT,

    (SELECT case
        when P.PARTY_AGE <= 18 then 'Underage'
        when P.PARTY_AGE between 19 and 21 then 'young 1'
        when P.PARTY_AGE between 22 and 24 then 'young 2'
        when P.PARTY_AGE between 24 and 60 then 'adult'
        when P.PARTY_AGE between 61 and 64 then 'elder 1'
        when P.PARTY_AGE >= 65 then 'elder 2' end as
age_range,
        COUNT(*) AS
TOTAL_NUMBER
    FROM PARTIES P
    WHERE P.PARTY_AGE IS NOT NULL
    group by (case
        when P.PARTY_AGE <= 18 then 'Underage'
        when P.PARTY_AGE between 19 and 21 then 'young
1'
        when P.PARTY_AGE between 22 and 24 then 'young
2'
        when P.PARTY_AGE between 24 and 60 then 'adult'
        when P.PARTY_AGE between 61 and 64 then 'elder
1'
        when P.PARTY_AGE >= 65 then 'elder 2'
    END)) TOTAL
WHERE TOTAL.age_range = FAULT.age_range
ORDER BY NUMBER_AT_FAULT / TOTAL_NUMBER DESC;

```

Query result (if the result is big, just a snippet)

| AGE_RANGE | RATIO_AT_FAULT |
|-----------|----------------|
| underage | 0.636 |
| young 1 | 0.572 |
| young 2 | 0.517 |
| elder 2 | 0.498 |
| adult | 0.409 |
| elder 1 | 0.399 |

Query 2:

Description of logic:

This query should retrieve the top 5 vehicles having the most collisions on roads with holes. We first do a subquery where we retrieve the ID of the vehicle type and the corresponding number of collisions. For that, we join the parties, road conditions and the relation between them, counting only the ones having holes, grouping by the ID of the vehicle type. We also sort it in descending order and fetch the 5 first row only in order to keep the 5 biggest values. We then use this subquery to extract the definition instead of the ID.

SQL statement

```
SELECT SWT.DEFINITION, STATS_COLLISIONS_HOLE.NUMBER_OF_COLLISION
FROM STATEWIDE_VEHICLE_TYPE SWT,
     (SELECT P.STATEWIDE_VEHICLE_TYPE_ID AS SVT_ID, COUNT(*) AS
NUMBER_OF_COLLISION
      FROM PARTIES P,
           COLLISION_WITH_ROAD_CONDITION CWRC,
           ROAD_CONDITION RC
     WHERE P.STATEWIDE_VEHICLE_TYPE_ID IS NOT NULL
           AND P.COLLISION_CASE_ID = CWRC.CASE_ID
           AND CWRC.ROAD_CONDITION_ID = RC.ID
           AND RC.DEFINITION = 'Holes, Deep Ruts'
     GROUP BY P.STATEWIDE_VEHICLE_TYPE_ID
     ORDER BY COUNT(*) DESC
     FETCH FIRST 5 ROW ONLY
    ) STATS_COLLISIONS_HOLE
WHERE SWT.ID = STATS_COLLISIONS_HOLE.SVT_ID;
```

Query result (if the result is big, just a snippet)

| DEFINITION | NUMBER_OF_COLLISIONS_HOLE |
|-------------------------------------|---------------------------|
| Passenger Car/Station Wagon | 10662 |
| Pickup or Panel Truck | 2263 |
| Motorcycle/Scooter | 450 |
| Bicycle | 430 |
| Truck or Truck Tractor with Trailer | 369 |

Query 3:

Description of logic:

This query should retrieve the top 10 vehicle makes with the most victims killed or with severe injuries. For this query, we first join the tables parties, victims and victim degree of injury. We then only take the rows where the degree of injury is either killed or severe injury and the vehicle make is not null (see our assumptions). We then group by the vehicle make, count the number of entries and sort it in descending order to be able to retrieve the top values only. Finally, we fetch the 10 first rows, in order to keep the top 10 only.

SQL statement

```
SELECT P.VEHICLE_MAKE, COUNT(*) AS  
NUMBER_OF_VICTIMS_KILLED_OR_WITH_SEVERE_INJURIES  
from PARTIES P,  
      VICTIMS V,  
      VICTIM_DEGREE_OF_INJURY VDOI  
WHERE P.ID = V.PARTY_ID  
      AND V.VICTIM_DEGREE_OF_INJURY_ID = VDOI.ID  
      AND (VDOI.DEFINITION = 'Killed' OR VDOI.DEFINITION = 'Severe  
Injury')  
      and P.VEHICLE_MAKE is not NULL -- NULL is the 4th more  
represented, not really interesting  
group by P.VEHICLE_MAKE  
order by COUNT(*) DESC  
      FETCH FIRST 10 ROW ONLY;
```

Query result (if the result is big, just a snippet)

| VEHICLE_MAKE | NUMBER_OF_VICTIMS_KILLED_OR_WITH_SEVERE_INJURIES |
|--------------|--|
| FORD | 13924 |
| HONDA | 12061 |
| TOYOTA | 10639 |
| CHEVROLET | 10418 |
| NISSAN | 3860 |
| DODGE | 3641 |

DIAS: Data-Intensive Applications and Systems Laboratory

School of Computer and Communication Sciences

Ecole Polytechnique Fédérale de Lausanne

Building BC, Station 14

CH-1015 Lausanne

URL: <http://dias.epfl.ch/>

| | |
|---------------------|------|
| HARLEY- DAVIDSON | 3410 |
| SUZUKI | 2482 |
| YAMAHA | 2105 |
| GMC | 1837 |

Query 4:

Description of logic:

This query should retrieve the safety index and the definition of the most safe and unsafe seating position. The safety factor is computed as the total number of victims having no injuries for a given position divided by the total number of victims at that position.

This query is done in two steps. We first create a table with every seating position and its corresponding safety factor. Then in a second query, we fetch only the row with the maximum and the row with the minimum safety factor.

For the creation of the table, we first run 2 subqueries. One that retrieves the definition of the position, and counts all the victims on that position. For that, we just join the tables for the victims and victims' degree of injuries and group by the seating position. We then extract the definition using the victim seating position table. The second subquery is almost equivalent except that we keep only the number of uninjured victims for each seating position. We then join the two subqueries on the definition of the seating position and compute the safety factor with a division of their respective count. To query the best and worst factors in this table, we then keep only the rows where the safety factor is either equivalent to the max or the min of the table, retrieved with 2 subqueries.

SQL statement

```
with SEATING_POSITION_TO_SAFETY_FACTOR AS (  
    SELECT UNINJURED.DEFINITION,  
           ROUND(UNINJURED.NUMBER_NO_INJURIES /  
ALL_DEGREES.NUMBER_ALL_DEGREE_INJURIES, 3) AS SAFETY_FACTOR  
    FROM (  
        SELECT VSP.DEFINITION,  
SEATING_POSITION_NO_INJURIES.NUMBER_NO_INJURIES as NUMBER_NO_INJURIES  
        FROM VICTIM_SEATING_POSITION VSP,  
        (  
            SELECT V.VICTIM_SEATING_POSITION_ID AS  
VICTIM_SEATING_POSITION_ID, COUNT(*) AS NUMBER_NO_INJURIES  
            FROM VICTIMS V,  
                VICTIM_DEGREE_OF_INJURY VDOI  
            WHERE V.VICTIM_DEGREE_OF_INJURY_ID = VDOI.ID  
                AND VDOI.DEFINITION = 'No Injury'  
                AND V.VICTIM_SEATING_POSITION_ID is not NULL  
            GROUP BY V.VICTIM_SEATING_POSITION_ID)  
SEATING_POSITION_NO_INJURIES  
        WHERE VSP.ID =  
SEATING_POSITION_NO_INJURIES.VICTIM_SEATING_POSITION_ID) UNINJURED,
```

```
(
    SELECT VSP.DEFINITION,
           GROUPED_SEATING_POSITIONS.NUMBER_ALL_DEGREE_INJURIES
as NUMBER_ALL_DEGREE_INJURIES
    FROM VICTIM_SEATING_POSITION VSP,
         (
            SELECT V.VICTIM_SEATING_POSITION_ID AS
VICTIM_SEATING_POSITION_ID,
            COUNT(*)                               AS
NUMBER_ALL_DEGREE_INJURIES
            FROM VICTIMS V,
                 VICTIM_DEGREE_OF_INJURY VDOI
            WHERE V.VICTIM_DEGREE_OF_INJURY_ID = VDOI.ID
                  AND VICTIM_SEATING_POSITION_ID is not NULL
            GROUP BY V.VICTIM_SEATING_POSITION_ID)
GROUPED_SEATING_POSITIONS
    WHERE VSP.ID =
GROUPED_SEATING_POSITIONS.VICTIM_SEATING_POSITION_ID) ALL_DEGREES

    WHERE UNINJURED.DEFINITION = ALL_DEGREES.DEFINITION)

SELECT *
FROM SEATING_POSITION_TO_SAFETY_FACTOR
WHERE SAFTEY_FACTOR = (SELECT MAX(SAFTEY_FACTOR) FROM
SEATING_POSITION_TO_SAFETY_FACTOR)
    OR SAFTEY_FACTOR = (SELECT MIN(SAFTEY_FACTOR) FROM
SEATING_POSITION_TO_SAFETY_FACTOR);
```

Query result (if the result is big, just a snippet)

| DEFINITION | SAFETY FACTOR |
|--------------------|---------------|
| DRIVER | 0.009 |
| STATION WAGON REAR | 0.825 |

Query 5:

Description of logic:

This query should retrieve the number of vehicle types which have had at least 10 collisions in at least half of the cities. For this query, we first keep the vehicles/locations tuples having at least 10 collisions. Join the tables parties and collisions, group by the vehicle type and city location (removing the null values) and count the number of entries for each of these tuples and keep only the ones having at least 10 entries. We then group by the type of vehicle and count the number of entries, which correspond to the number of cities in which each vehicle type had at least 10 collisions. We then only keep those where this number is at least half of the cities. To count half the number of the cities, we used a subquery which counts every unique location and divides it by 2.

SQL statement

```
SELECT COUNT(*) AS NUMBER_OF_VEHICLE_TYPE
FROM (SELECT TYPE_CITY_TO_ACCIDENT_COUNT.TYPE
      FROM (SELECT P.STATEWIDE_VEHICLE_TYPE_ID AS TYPE,
                  C.COUNTY_CITY_LOCATION
              FROM PARTIES P,
                  COLLISIONS C
              WHERE P.COLLISION_CASE_ID = C.CASE_ID
                   AND C.COUNTY_CITY_LOCATION IS NOT NULL
                   AND P.STATEWIDE_VEHICLE_TYPE_ID IS NOT NULL
              GROUP BY (P.STATEWIDE_VEHICLE_TYPE_ID,
                       C.COUNTY_CITY_LOCATION)
              HAVING COUNT(*) >= 10
            ) TYPE_CITY_TO_ACCIDENT_COUNT
      GROUP BY TYPE_CITY_TO_ACCIDENT_COUNT.TYPE
      HAVING COUNT(*) >= (SELECT COUNT(UNIQUE
                                (C.COUNTY_CITY_LOCATION)) / 2
                          FROM COLLISIONS C
                        )
    ) TYPE_TO_CITY_COUNT;
```

Query result (if the result is big, just a snippet)

| NUMBER_OF_VEHICLE_TYPE |
|------------------------|
|------------------------|

| |
|----|
| 13 |
|----|

Query 6:

Description of logic:

This query should retrieve the top 10 minimum average age cases for the 3 most populated cities, together with the city location, the population and the case id.

In order to do that, we rely on some subqueries. First, we compute the average victim age for each collision that happened in the 3 most populated cities. To get those 3 cities, we simply take 3 cities (no specific ordering) that have a population_id that corresponds to 'Incorporated (over 250000)'. Once we have the average victim age for each collision in the top 3 most populated cities, we label each resulting average victim age in ascending order. This will label the average victim age in each city. This allows us to start the count at 1 for each of the 3 cities. The last part of the query consists of taking the resulting rows that have a row_number less or equal to 10. This way we can show for each of the top-3 most populated cities the bottom 10 collisions in terms of average victim age.

Remark:

- We only showed the first 20 entries in the result as asked in the question.
- Due to our assumptions on the null values (see Assumptions), we only got 0 as age average. We would probably have had some non-zero values if we had discarded the cases where some ages were null, but as stated in the assumption, it made sense to us to discard as little data as possible.

SQL statement

```
with average_age(COLLISION_CASE_ID, COUNTY_CITY_LOCATION,
POPULATION_ID, V_AGE) as
(
    SELECT distinct COLLISION_CASE_ID,
                   COUNTY_CITY_LOCATION,
                   POPULATION_ID,
                   avg(v.VICTIM_AGE) OVER (PARTITION BY
C.CASE_ID) as v_age
    FROM COLLISIONS C
         INNER JOIN PARTIES on C.CASE_ID =
PARTIES.COLLISION_CASE_ID
         inner join VICTIMS V on PARTIES.ID =
V.PARTY_ID
    WHERE C.COUNTY_CITY_LOCATION in (
        SELECT distinct COUNTY_CITY_LOCATION
        from COLLISIONS C
        INNER JOIN POPULATION P ON P.ID =
C.POPULATION_ID
        where C.POPULATION_ID in
```



```
(
    SELECT distinct (C.POPULATION_ID)
    FROM COLLISIONS C
    WHERE P.DEFINITION = 'Incorporated (over
250000) '
)
    FETCH FIRST 3 ROWS ONLY
),
rws as (
    SELECT ROW_NUMBER() OVER (PARTITION BY
COUNTY_CITY_LOCATION
    ORDER BY V_AGE ASC ) AS Row_Number,
    COLLISION_CASE_ID,
    COUNTY_CITY_LOCATION,
    POPULATION_ID,
    V_AGE
    FROM average_age
)
select COLLISION_CASE_ID, COUNTY_CITY_LOCATION, P.DEFINITION,
V_AGE as AVERAGE_VICTIM_AGE
from rws
    INNER JOIN POPULATION P ON P.ID = POPULATION_ID
where Row_Number <= 10
order by COUNTY_CITY_LOCATION, V_AGE asc;
```

Query result (if the result is big, just a snippet)

| COLLISION_CASE_ID | COUNTY_CITY_LOCATION | DEFINITION | AVERAGE_VICTIM_AGE |
|-------------------|----------------------|----------------------------------|--------------------|
| 1838702 | 109 | Incorporated (over 250000) | 0 |
| 2727453 | 109 | Incorporated (over 250000) | 0 |
| 3486455 | 109 | Incorporated (over 250000) | 0 |

DIAS: Data-Intensive Applications and Systems Laboratory

School of Computer and Communication Sciences

Ecole Polytechnique Fédérale de Lausanne

Building BC, Station 14

CH-1015 Lausanne

URL: <http://dias.epfl.ch/>

| | | | |
|---------|------|----------------------------------|---|
| 0059033 | 109 | Incorporated (over 250000) | 0 |
| 2295152 | 109 | Incorporated (over 250000) | 0 |
| 1336621 | 109 | Incorporated (over 250000) | 0 |
| 1231119 | 109 | Incorporated (over 250000) | 0 |
| 2737180 | 109 | Incorporated (over 250000) | 0 |
| 2506007 | 109 | Incorporated (over 250000) | 0 |
| 1377820 | 109 | Incorporated (over 250000) | 0 |
| 2715062 | 3019 | Incorporated (over 250000) | 0 |
| 2412373 | 3019 | Incorporated (over 250000) | 0 |
| 1994820 | 3019 | Incorporated (over 250000) | 0 |
| 1170908 | 3019 | Incorporated (over 250000) | 0 |

DIAS: Data-Intensive Applications and Systems Laboratory

School of Computer and Communication Sciences

Ecole Polytechnique Fédérale de Lausanne

Building BC, Station 14

CH-1015 Lausanne

URL: <http://dias.epfl.ch/>

| | | | |
|---------|------|----------------------------------|---|
| 1825689 | 3019 | Incorporated (over 250000) | 0 |
| 3553649 | 3019 | Incorporated (over 250000) | 0 |
| 2072101 | 3019 | Incorporated (over 250000) | 0 |
| 2138547 | 3019 | Incorporated (over 250000) | 0 |
| 2674015 | 3019 | Incorporated (over 250000) | 0 |
| 3551315 | 3019 | Incorporated (over 250000) | 0 |

Query 7:

Description of logic:

This query should retrieve all the collisions of type pedestrian where all the victims were above 100 years old. We should then show only the `collision_id` and the `victim_age` of the oldest victim for each of them.

For this query, we first joined the 4 tables used (victims, parties, collisions and type of collisions) and kept only those that are of type pedestrian. We then grouped by the `case_id` and kept only the collisions where the minimum `victim_age` is above 100, to be sure that all victims were older than 100. We then returned the `case_id` and maximum `victim_age`.

Remark:

We only showed the first 20 entries in the result as asked in the question.

SQL statement

```
SELECT C.CASE_ID, MAX(V.VICTIM_AGE) AS AGE_MAX
FROM VICTIMS V, PARTIES P, COLLISIONS C, TYPE_OF_COLLISION TOC
WHERE V.PARTY_ID = P.ID
  AND P.COLLISION_CASE_ID = C.CASE_ID
  AND C.TYPE_OF_COLLISION_ID = TOC.ID
  AND TOC.DEFINITION = 'Vehicle/Pedestrian'
GROUP BY CASE_ID
HAVING MIN(V.VICTIM_AGE) > 100;
```

Query result (if the result is big, just a snippet)

| CASE_ID | AGE_MAX |
|---------|---------|
| 2531557 | 103 |
| 0439197 | 102 |
| 1548445 | 102 |
| 1373664 | 101 |
| 1209166 | 101 |
| 1347636 | 101 |
| 0828116 | 102 |
| 0784061 | 102 |
| 1213340 | 121 |
| 0817210 | 102 |
| 0036446 | 110 |
| 3485436 | 101 |
| 0820619 | 101 |
| 0868472 | 103 |
| 1847678 | 104 |
| 0644226 | 103 |
| 0566220 | 102 |
| 3388544 | 105 |
| 2472739 | 103 |
| 0851026 | 106 |

Query 8:

Description of logic:

This query should retrieve the vehicles which have participated in at least 10 collisions and their corresponding number of accidents. For the vehicle ID, see our assumptions.

For this query, we first join the parties and the vehicle types (to retrieve the definition). We then keep only the vehicles having not any null values (for the make, the year and the type) and group them together. We count them and only keep those that appear at least 10 times. We finally sort them in descending order.

We can observe that the type of the vehicles having the most collisions is always “Passenger Car/Station Wagon” which is quite logical since it represents the most common type. We can also observe that the make is always either TOYOTA, FORD or HONDA and the year between 1997 and 2002 for the top 20 vehicles in terms of collisions.

Remark:

We only showed the first 20 entries in the result as asked in the question.

SQL statement

```
SELECT SVT.DEFINITION, P.VEHICLE_MAKE, P.VEHICLE_YEAR, COUNT(*) AS  
NUMBER_COLLISION  
FROM PARTIES P,  
STATEWIDE_VEHICLE_TYPE SVT  
WHERE P.STATEWIDE_VEHICLE_TYPE_ID IS NOT NULL  
AND P.VEHICLE_MAKE IS NOT NULL  
AND P.VEHICLE_YEAR IS NOT NULL  
AND P.STATEWIDE_VEHICLE_TYPE_ID = SVT.ID  
GROUP BY (SVT.DEFINITION, P.VEHICLE_MAKE, P.VEHICLE_YEAR)  
HAVING COUNT(*) >= 10  
ORDER BY COUNT(*) DESC;
```

Query result (if the result is big, just a snippet)

| DEFINITION | VEHICLE_MAKE | VEHICLE_YEAR | NUMBER_COLLISION |
|-----------------------------|--------------|--------------|------------------|
| Passenger Car/Station Wagon | TOYOTA | 2000 | 52504 |
| Passenger Car/Station Wagon | FORD | 2000 | 51943 |
| Passenger Car/Station Wagon | HONDA | 2000 | 50284 |
| Passenger Car/Station Wagon | FORD | 1998 | 49182 |
| Passenger Car/Station Wagon | TOYOTA | 2001 | 47232 |
| Passenger Car/Station Wagon | HONDA | 2001 | 45277 |
| Passenger Car/Station Wagon | FORD | 2001 | 45236 |
| Passenger Car/Station Wagon | TOYOTA | 1999 | 42941 |
| Passenger Car/Station Wagon | HONDA | 1998 | 42091 |
| Passenger Car/Station Wagon | FORD | 1999 | 41948 |
| Passenger Car/Station Wagon | FORD | 1995 | 40246 |
| Passenger Car/Station Wagon | HONDA | 1997 | 39210 |
| Passenger Car/Station Wagon | FORD | 1997 | 38885 |
| Passenger Car/Station Wagon | HONDA | 1999 | 38556 |

DIAS: Data-Intensive Applications and Systems Laboratory

School of Computer and Communication Sciences

Ecole Polytechnique Fédérale de Lausanne

Building BC, Station 14

CH-1015 Lausanne

URL: <http://dias.epfl.ch/>

| | | | |
|-----------------------------|--------|------|-------|
| Passenger Car/Station Wagon | TOYOTA | 2002 | 38427 |
| Passenger Car/Station Wagon | TOYOTA | 1998 | 38012 |
| Passenger Car/Station Wagon | TOYOTA | 1997 | 37158 |
| Passenger Car/Station Wagon | TOYOTA | 2003 | 35943 |
| Passenger Car/Station Wagon | HONDA | 2002 | 35785 |
| Passenger Car/Station Wagon | FORD | 2002 | 35460 |

Query 9:

Description of logic:

This query should retrieve the top 10 cities having the most collisions.

For that, we simply group by the county city location, count the number of entries. To retrieve the top 10 cities only, we sort our result in descending order and fetch the first 10 rows only.

SQL statement

```
SELECT COUNTY_CITY_LOCATION, COUNT(*) AS NUMBER_COLLISIONS
FROM COLLISIONS C
GROUP BY COUNTY_CITY_LOCATION
ORDER BY NUMBER_COLLISIONS DESC FETCH FIRST 10 ROWS ONLY;
```

Query result (if the result is big, just a snippet)

| COUNTY_CITY_LOCATION | NUMBER_COLLISIONS |
|----------------------|-------------------|
| 1942 | 399582 |
| 1900 | 118446 |
| 3400 | 80191 |
| 3711 | 76867 |
| 109 | 72995 |
| 3300 | 61453 |
| 3404 | 58068 |
| 4313 | 57852 |
| 1941 | 53565 |
| 3801 | 48450 |

Query 10:

Description of logic:

This query should retrieve the number of accidents for different time periods with different lighting conditions.

For this query, we decided to first take into account the lighting conditions that were not ambiguous, i.e. daylight for day and everything containing dark for night. For the dusk/dawn category, we looked at the time and month and the information given in the question to put them in the right category. We decided to discard all the data that was not consistent (for example if an accident had lighting 'Dusk-Dawn' but was at time 12:00 which is neither dusk nor dawn, we dropped it).

If the lighting condition was not given (null), we tried to infer the period based on the time only when it was possible or the time and the date when both were available.

SQL statement

```
SELECT TIME_PERIOD, COUNT(*) as NUMBER_ACCIDENT
FROM (
    SELECT CASE
        when l.DEFINITION = 'Daylight' then 'DAY_COLLISIONS'
        when l.DEFINITION like '%dark%' then 'NIGHT_COLLISIONS'
        when l.DEFINITION = 'Dusk - Dawn' then
            case
                when C.COLLISION_DATE is not null then
                    case
                        WHEN ((EXTRACT(MONTH FROM C.COLLISION_DATE)
BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '20' AND '21'))
                        OR (EXTRACT(MONTH FROM
C.COLLISION_DATE) NOT BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '18' AND '19'))
                        THEN 'DUSK_COLLISIONS'
                        WHEN ((EXTRACT(MONTH FROM C.COLLISION_DATE)
BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '4' AND '5'))
                        OR (EXTRACT(MONTH FROM
C.COLLISION_DATE) NOT BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '6' AND '7'))
                        THEN 'DAWN_COLLISIONS'
```

```

        end
    end

    else
        case
            when C.COLLISION_DATE is not null then
                CASE
                    WHEN ((EXTRACT(MONTH FROM C.COLLISION_DATE)
BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '20' AND '21'))
                        OR (EXTRACT(MONTH FROM
C.COLLISION_DATE) NOT BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '18' AND '19'))
                            THEN 'DUSK_COLLISIONS'
                    WHEN ((EXTRACT(MONTH FROM C.COLLISION_DATE)
BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '4' AND '5'))
                        OR (EXTRACT(MONTH FROM
C.COLLISION_DATE) NOT BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '6' AND '7'))
                            THEN 'DAWN_COLLISIONS'
                    WHEN (EXTRACT(MONTH FROM C.COLLISION_DATE)
BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '6' AND '19'))
                        OR (EXTRACT(MONTH FROM
C.COLLISION_DATE) NOT BETWEEN '4' AND '8'
                                AND EXTRACT(HOUR FROM
C.COLLISION_TIME) BETWEEN '8' AND '17'))
                            THEN 'DAY_COLLISIONS'
                    ELSE 'NIGHT_COLLISIONS'
                end
            end
        else
            case
                when extract(hour from C.COLLISION_TIME) > 7
                    and extract(hour from C.COLLISION_TIME) <
18 then 'DAY_COLLISIONS'
                when extract(hour from C.COLLISION_TIME) < 4
                    and extract(hour from C.COLLISION_TIME) >
21 then 'NIGHT_COLLISIONS'
            end
        end
    end
end
```

```
        end
    end as TIME_PERIOD
FROM COLLISIONS C
left outer join LIGHTING L on C.LIGHTING_ID = L.ID
)
where TIME_PERIOD is not null
GROUP BY TIME_PERIOD
ORDER BY NUMBER_ACCIDENT DESC;
```

Query result (if the result is big, just a snippet)

| TIME_PERIOD | NUMBER_ACCIDENT |
|------------------|-----------------|
| DAY_COLLISIONS | 2607362 |
| NIGHT_COLLISIONS | 628870 |
| DUSK_COLLISIONS | 305720 |
| DAWN_COLLISIONS | 64534 |

Query Performance Analysis – Indexing

We observed that the running time for a query varies for each run. Therefore we took the mean of five runs for the initial and optimized time. However, as time can vary a lot between different runs, we tried to rather use the cost of the plan to gain more insightful information about the quality of our optimization.

Query 1

Initial time: 4,5s

Optimized time: 2,5s

Explain the improvement:

We created the following 2 indexes:

```
CREATE INDEX PARTIES_IDX_PARTY_AGE on PARTIES (PARTY_AGE);
CREATE INDEX PARTIES_IDX_AT_FAULT_PARTY_AGE on PARTIES (AT_FAULT,
PARTY_AGE);
```

In this query we only access the parties table, once over the `party_age` and once over the `party_age` and the `at_fault` attributes. So it makes perfect sense to create one index for each of these accesses and we can indeed see on the optimized plan that both the `TABLE ACCESS FULL` have been replaced by `INDEX FAST FULL SCAN`. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

Initial plan:

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-------------------|---------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 112 | 4256 | 60028 (1) | 00:00:03 |
| 1 | SORT ORDER BY | | 112 | 4256 | 60028 (1) | 00:00:03 |
| * 2 | HASH JOIN | | 112 | 4256 | 60027 (1) | 00:00:03 |
| 3 | VIEW | | 106 | 2014 | 29997 (1) | 00:00:02 |
| 4 | HASH GROUP BY | | 106 | 530 | 29997 (1) | 00:00:02 |
| * 5 | TABLE ACCESS FULL | PARTIES | 2808K | 13M | 29927 (1) | 00:00:02 |
| 6 | VIEW | | 106 | 2014 | 30030 (1) | 00:00:02 |
| 7 | HASH GROUP BY | | 106 | 318 | 30030 (1) | 00:00:02 |
| * 8 | TABLE ACCESS FULL | PARTIES | 6188K | 17M | 29868 (1) | 00:00:02 |

Predicate Information (identified by operation id):

```
" 2 - access("TOTAL"."AGE_RANGE"="FAULT"."AGE_RANGE") "
```

```
" 5 - filter("P"."PARTY_AGE" IS NOT NULL AND "P"."AT_FAULT"='T') "
```

```
" 8 - filter("P"."PARTY_AGE" IS NOT NULL) "
```

Improved plan:

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|----------------------|--------------------------------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 112 | 4256 | 7926 (5) | 00:00:01 |
| 1 | SORT ORDER BY | | 112 | 4256 | 7926 (5) | 00:00:01 |
| * 2 | HASH JOIN | | 112 | 4256 | 7925 (5) | 00:00:01 |
| 3 | VIEW | | 106 | 2014 | 4463 (4) | 00:00:01 |
| 4 | HASH GROUP BY | | 106 | 530 | 4463 (4) | 00:00:01 |
| * 5 | INDEX FAST FULL SCAN | PARTIES_IDX_AT_FAULT_PARTY_AGE | 2808K | 13M | 4393 (2) | 00:00:01 |
| 6 | VIEW | | 106 | 2014 | 3461 (6) | 00:00:01 |
| 7 | HASH GROUP BY | | 106 | 318 | 3461 (6) | 00:00:01 |
| * 8 | INDEX FAST FULL SCAN | PARTIES_IDX_PARTY_AGE | 6188K | 17M | 3299 (1) | 00:00:01 |

Predicate Information (identified by operation id):

```

" 2 - access("TOTAL"."AGE_RANGE"="FAULT"."AGE_RANGE")
" 5 - filter("P"."PARTY_AGE" IS NOT NULL AND "P"."AT_FAULT"='T')
" 8 - filter("P"."PARTY_AGE" IS NOT NULL)

```

Query 2

Initial time: 1,5s

Optimized time: 2,2s

Explain the improvement:

We created the following 3 indexes:

```

CREATE INDEX
PARTIES_IDX_COLLISION_CASE_ID_STATEWIDE_VEHICLE_TYPE_ID on
PARTIES (STATEWIDE_VEHICLE_TYPE_ID, COLLISION_CASE_ID);
CREATE INDEX ROAD_CONDITION_IDX_DEFINITION_ID on
ROAD_CONDITION (DEFINITION, ID);
CREATE INDEX STATEWIDE_VEHICLE_TYPE_IDX_DEFINITION_ID on
STATEWIDE_VEHICLE_TYPE (DEFINITION, ID);

```

We created an index for every group of attributes per accessed table. An index on COLLISION_WITH_ROAD_CONDITION was not needed since the accessed tuple is the primary key of the table so it's already clustered on that index. Except for this case all the other TABLE ACCESS FULL, on the improved plan have been replaced by INDEX FAST FULL SCAN. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

Initial plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|------|-----------------------------|-------------------------------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 4 | 252 | | 65994 (1) | 00:00:03 |
| 1 | SORT ORDER BY | | 4 | 252 | | 65994 (1) | 00:00:03 |
| 2 | MERGE JOIN | | 4 | 252 | | 65993 (1) | 00:00:03 |
| 3 | TABLE ACCESS BY INDEX ROWID | STATEWIDE_VEHICLE_TYPE | 15 | 330 | | 2 (0) | 00:00:01 |
| 4 | INDEX FULL SCAN | SYS_C00207107 | 15 | | | 1 (0) | 00:00:01 |
| * 5 | SORT JOIN | | 5 | 205 | | 65991 (1) | 00:00:03 |
| * 6 | VIEW | | 5 | 205 | | 65990 (1) | 00:00:03 |
| * 7 | WINDOW SORT PUSHED RANK | | 15 | 2325 | | 65990 (1) | 00:00:03 |
| 8 | HASH GROUP BY | | 15 | 2325 | | 65990 (1) | 00:00:03 |
| * 9 | HASH JOIN | | 806K | 119M | 43M | 65951 (1) | 00:00:03 |
| * 10 | HASH JOIN | | 456K | 38M | | 9954 (1) | 00:00:01 |
| * 11 | TABLE ACCESS FULL | ROAD_CONDITION | 1 | 21 | | 3 (0) | 00:00:01 |
| 12 | TABLE ACCESS FULL | COLLISION_WITH_ROAD_CONDITION | 3652K | 233M | | 9942 (1) | 00:00:01 |
| * 13 | TABLE ACCESS FULL | PARTIES | 6400K | 408M | | 29906 (1) | 00:00:02 |

Predicate Information (identified by operation id):

```

5 - access("SWT"."ID"="from$_subquery$_006"."SVT_ID")
   filter("SWT"."ID"="from$_subquery$_006"."SVT_ID")
6 - filter("from$_subquery$_006"."rowlimit $$ rownumber"<=5)
7 - filter(ROW_NUMBER() OVER ( ORDER BY COUNT(*) DESC )<=5)
9 - access("P"."COLLISION_CASE_ID"="CWRC"."CASE_ID")
10 - access("CWRC"."ROAD_CONDITION_ID"="RC"."ID")
11 - filter("RC"."DEFINITION"='Holes, Deep Ruts')
13 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL)

```

Improved plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|------|-------------------------|---|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 4 | 252 | | 57549 (1) | 00:00:03 |
| 1 | SORT ORDER BY | | 4 | 252 | | 57549 (1) | 00:00:03 |
| * 2 | HASH JOIN | | 4 | 252 | | 57548 (1) | 00:00:03 |
| * 3 | VIEW | | 5 | 205 | | 57547 (1) | 00:00:03 |
| * 4 | WINDOW SORT PUSHED RANK | | 15 | 2325 | | 57547 (1) | 00:00:03 |
| 5 | HASH GROUP BY | | 15 | 2325 | | 57547 (1) | 00:00:03 |
| * 6 | HASH JOIN | | 806K | 119M | 43M | 57508 (1) | 00:00:03 |
| * 7 | HASH JOIN | | 456K | 38M | | 9952 (1) | 00:00:01 |
| * 8 | INDEX RANGE SCAN | ROAD_CONDITION_IDX_DEFINITION_ID | 1 | 21 | | 1 (0) | 00:00:01 |
| 9 | TABLE ACCESS FULL | COLLISION_WITH_ROAD_CONDITION | 3652K | 233M | | 9942 (1) | 00:00:01 |
| * 10 | INDEX FAST FULL SCAN | PARTIES_IDX_COLLISION_CASE_ID_STATEWIDE_VEHICLE_TYPE_ID | 6400K | 408M | | 21465 (1) | 00:00:01 |
| 11 | INDEX FULL SCAN | STATEWIDE_VEHICLE_TYPE_IDX_DEFINITION_ID | 15 | 330 | | 1 (0) | 00:00:01 |

Predicate Information (identified by operation id):

```

2 - access("SWT"."ID"="from$_subquery$_006"."SVT_ID")
3 - filter("from$_subquery$_006"."rowlimit $$ rownumber"<=5)
4 - filter(ROW_NUMBER() OVER ( ORDER BY COUNT(*) DESC )<=5)
6 - access("P"."COLLISION_CASE_ID"="CWRC"."CASE_ID")
7 - access("CWRC"."ROAD_CONDITION_ID"="RC"."ID")
8 - access("RC"."DEFINITION"='Holes, Deep Ruts')
10 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL)

```

Query 3

Initial time: 4s

Optimized time: 1,7s

Explain the improvement:

We created the following 4 indexes:

```
CREATE INDEX PARTIES_IDX_VEHICLE_MAKE on PARTIES(VEHICLE_MAKE);
CREATE INDEX PARTIES_IDX_ID_VEHICLE_MAKE on PARTIES(VEHICLE_MAKE,
ID);
CREATE INDEX VICTIMS_IDX_PARTY_ID_VICTIM_DEGREE_OF_INJURY_ID on
VICTIMS(VICTIM_DEGREE_OF_INJURY_ID, PARTY_ID);
CREATE INDEX VICTIM_DEGREE_OF_INJURY_IDX_DEFINITION_ID on
VICTIM_DEGREE_OF_INJURY(DEFINITION, ID);
```

In this query, we access the party table once over the `vehicle_make` and the `id` and therefore our second index on those attributes improves the plan by replacing the `TABLE ACCESS FULL` with an `INDEX FAST FULL SCAN`. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

The same improvement happens for the table `victims` and the attributes `VICTIM_DEGREE_OF_INJURY_ID` and `PARTY_ID` thanks to our third index. The victim degree of injury table is also accessed over its `id` and `definition` and therefore it made sense to use an index, which transforms the `TABLE ACCESS FULL` into an `INDEX RANGE SCAN`. Our first index doesn't change the plan, but it improved the cost a bit, probably because it is used during the group by.

Initial plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|----|-------------------------|-------------------------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 10 | 1160 | | 45925 (1) | 00:00:02 |
| 1 | SORT ORDER BY | | 10 | 1160 | | 45925 (1) | 00:00:02 |
| 2 | VIEW | | 10 | 1160 | | 45924 (1) | 00:00:02 |
| 3 | WINDOW SORT PUSHED RANK | | 209 | 8778 | | 45924 (1) | 00:00:02 |
| 4 | HASH GROUP BY | | 209 | 8778 | | 45924 (1) | 00:00:02 |
| 5 | HASH JOIN | | 1360K | 54M | 53M | 45858 (1) | 00:00:02 |
| 6 | HASH JOIN | | 1360K | 37M | | 5283 (2) | 00:00:01 |
| 7 | TABLE ACCESS FULL | VICTIM_DEGREE_OF_INJURY | 2 | 40 | | 3 (0) | 00:00:01 |
| 8 | TABLE ACCESS FULL | VICTIMS | 4082K | 35M | | 5269 (1) | 00:00:01 |
| 9 | TABLE ACCESS FULL | PARTIES | 6759K | 83M | | 29909 (1) | 00:00:02 |

Predicate Information (identified by operation id):


```
" 2 - filter("from$ subquery$ _004"."rowlimit $$ rownumber"<=10) "
3 - filter(ROW_NUMBER() OVER ( ORDER BY COUNT(*) DESC )<=10)
" 5 - access("P"."ID"="V"."PARTY_ID") "
" 6 - access("V"."VICTIM_DEGREE_OF_INJURY_ID"="VDOI"."ID") "
" 7 - filter("VDOI"."DEFINITION"='Killed' OR "VDOI"."DEFINITION"='Severe Injury') "
" 8 - filter("V"."VICTIM_DEGREE_OF_INJURY_ID">=0 AND "V"."VICTIM_DEGREE_OF_INJURY_ID"<=7) "
" 9 - filter("P"."VEHICLE_MAKE" IS NOT NULL) "
```

Improved plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost | (%CPU) | Time |
|----|-------------------------|---|-------|-------|---------|-------|--------|----------|
| 0 | SELECT STATEMENT | | 10 | 1160 | | 20343 | (1) | 00:00:01 |
| 1 | SORT ORDER BY | | 10 | 1160 | | 20343 | (1) | 00:00:01 |
| 2 | VIEW | | 10 | 1160 | | 20342 | (1) | 00:00:01 |
| 3 | WINDOW SORT PUSHED RANK | | 209 | 8778 | | 20342 | (1) | 00:00:01 |
| 4 | HASH GROUP BY | | 209 | 8778 | | 20342 | (1) | 00:00:01 |
| 5 | HASH JOIN | | 1360K | 54M | 53M | 20276 | (1) | 00:00:01 |
| 6 | HASH JOIN | | 1360K | 37M | | 3021 | (2) | 00:00:01 |
| 7 | INLIST ITERATOR | | | | | | | |
| 8 | INDEX RANGE SCAN | VICTIM_DEGREE_OF_INJURY_IDX_DEFINITION_ID | 2 | 40 | | 1 | (0) | 00:00:01 |
| 9 | INDEX FAST FULL SCAN | VICTIMS_IDX_PARTY_ID_VICTIM_DEGREE_OF_INJURY_ID | 4082K | 35M | | 3010 | (1) | 00:00:01 |
| 10 | INDEX FAST FULL SCAN | PARTIES_IDX_ID_VEHICLE_MAKE | 6759K | 83M | | 6589 | (1) | 00:00:01 |

Predicate Information (identified by operation id):

```
" 2 - filter("from$ subquery$ _004"."rowlimit $$ rownumber"<=10) "
3 - filter(ROW_NUMBER() OVER ( ORDER BY COUNT(*) DESC )<=10)
" 5 - access("P"."ID"="V"."PARTY_ID") "
" 6 - access("V"."VICTIM_DEGREE_OF_INJURY_ID"="VDOI"."ID") "
" 8 - access("VDOI"."DEFINITION"='Killed' OR "VDOI"."DEFINITION"='Severe Injury') "
" 9 - filter("V"."VICTIM_DEGREE_OF_INJURY_ID">=0 AND "V"."VICTIM_DEGREE_OF_INJURY_ID"<=7) "
" 10 - filter("P"."VEHICLE_MAKE" IS NOT NULL) "
```

Query 5

Initial time: 3 min 20 s

Optimized time: 55s

Explain the improvement:

We created the following 2 indexes:

```
CREATE INDEX
PARTIES_IDX_COLLISION_CASE_ID_STATEWIDE_VEHICLE_TYPE_ID on
PARTIES (STATEWIDE_VEHICLE_TYPE_ID, COLLISION_CASE_ID);
CREATE INDEX COLLISIONS_IDX_CASE_ID_COUNTY_CITY_LOCATION on
COLLISIONS (COUNTY_CITY_LOCATION, CASE_ID);
```

By creating an index on (STATEWIDE_VEHICLE_TYPE_ID, COLLISION_CASE_ID) and on (COUNTY_CITY_LOCATION, CASE_ID), we are able to replace all the TABLE ACCESS FULL by INDEX FAST FULL SCAN. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

Initial plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|------|-------------------|------------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | | | 87192 (1) | 00:00:04 |
| 1 | SORT AGGREGATE | | 1 | | | | |
| 2 | VIEW | | 1 | | | 87192 (1) | 00:00:04 |
| * 3 | FILTER | | | | | | |
| 4 | HASH GROUP BY | | 1 | 2 | | 87192 (1) | 00:00:04 |
| 5 | VIEW | | 287 | 574 | | 87192 (1) | 00:00:04 |
| * 6 | FILTER | | | | | | |
| 7 | HASH GROUP BY | | 287 | 39032 | | 87192 (1) | 00:00:04 |
| * 8 | HASH JOIN | | 6400K | 830M | 284M | 87024 (1) | 00:00:04 |
| * 9 | TABLE ACCESS FULL | COLLISIONS | 3678K | 242M | | 19090 (1) | 00:00:01 |
| * 10 | TABLE ACCESS FULL | PARTIES | 6400K | 408M | | 29906 (1) | 00:00:02 |
| 11 | SORT AGGREGATE | | 1 | 13 | | | |
| 12 | VIEW | VM_NWVW_1 | 540 | 7020 | | 19182 (1) | 00:00:01 |
| 13 | SORT GROUP BY | | 540 | 2160 | | 19182 (1) | 00:00:01 |
| 14 | TABLE ACCESS FULL | COLLISIONS | 3678K | 14M | | 19088 (1) | 00:00:01 |

Predicate Information (identified by operation id):

```

3 - filter(COUNT(*)>= (SELECT COUNT("$vm_col_1")/2 FROM (SELECT "
      "C"."COUNTY_CITY_LOCATION" "$vm_col_1" FROM "COLLISIONS" "C" GROUP BY "
      "C"."COUNTY_CITY_LOCATION") "$vm_nwvw_1"))
6 - filter(COUNT(*)>=10)
8 - access("P"."COLLISION_CASE_ID"="C"."CASE_ID")
9 - filter("C"."COUNTY_CITY_LOCATION" IS NOT NULL)
10 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL)

```

Improved plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|------|----------------------|---|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 1 | | | 70847 (1) | 00:00:03 |
| 1 | SORT AGGREGATE | | 1 | | | | |
| 2 | VIEW | | 1 | | | 70847 (1) | 00:00:03 |
| * 3 | FILTER | | | | | | |
| 4 | HASH GROUP BY | | 1 | 2 | | 70847 (1) | 00:00:03 |
| 5 | VIEW | | 287 | 574 | | 70847 (1) | 00:00:03 |
| * 6 | FILTER | | | | | | |
| 7 | HASH GROUP BY | | 287 | 39032 | | 70847 (1) | 00:00:03 |
| * 8 | HASH JOIN | | 6400K | 830M | 284M | 70679 (1) | 00:00:03 |
| * 9 | INDEX FAST FULL SCAN | COLLISIONS_IDX_CASE_ID_COUNTY_CITY_LOCATION | 3678K | 242M | | 11187 (1) | 00:00:01 |
| * 10 | INDEX FAST FULL SCAN | PARTIES_IDX_COLLISION_CASE_ID_STATEWIDE_VEHICLE_TYPE_ID | 6400K | 408M | | 21465 (1) | 00:00:01 |
| 11 | SORT AGGREGATE | | 1 | 13 | | | |
| 12 | VIEW | VM_NWVW_1 | 540 | 7020 | | 11278 (1) | 00:00:01 |
| 13 | SORT GROUP BY | | 540 | 2160 | | 11278 (1) | 00:00:01 |
| 14 | INDEX FAST FULL SCAN | COLLISIONS_IDX_CASE_ID_COUNTY_CITY_LOCATION | 3678K | 14M | | 11185 (1) | 00:00:01 |

Predicate Information (identified by operation id):

```

3 - filter(COUNT(*)>= (SELECT COUNT("$vm_col_1")/2 FROM (SELECT "C"."COUNTY_CITY_LOCATION" "$vm_col_1" FROM "COLLISIONS" "C" GROUP BY "C"."COUNTY_CITY_LOCATION") "$vm_nwvw_1"))
6 - filter(COUNT(*)>=10)
8 - access("P"."COLLISION_CASE_ID"="C"."CASE_ID")
9 - filter("C"."COUNTY_CITY_LOCATION" IS NOT NULL)
10 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL)

```

Query 7

Initial time: 23,47s

Optimized time: 8,96s

Explain the improvement:

We created the following 4 indexes:

```
CREATE INDEX VICTIMS_IDX_PARTY_ID_VICTIM_AGE on
VICTIMS (VICTIM_AGE, PARTY_ID);
CREATE INDEX PARTIES_IDX_COLLISION_CASE_ID_ID on
PARTIES (COLLISION_CASE_ID, ID);
CREATE INDEX COLLISIONS_IDX_CASE_ID_TYPE_OF_COLLISION_ID on
COLLISIONS (TYPE_OF_COLLISION_ID, CASE_ID);
CREATE INDEX TYPE_OF_COLLISION_IDX_DEFINITION_ID on
TYPE_OF_COLLISION (DEFINITION, ID);
```

In the query we access the three main tables so creating an index for each of these accesses greatly improves the runtime. We also created an index on (TYPE_OF_COLLISION) since it's a rather small table it doesn't reduce the cost much, but it's still an improvement. By creating an index on parties we avoid a TABLE ACCESS FULL and replace it by an INDEX FAST FULL SCAN. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

Initial plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost (%CPU) | Time |
|-----|-------------------|-------------------|-------|-------|---------|-------------|----------|
| 0 | SELECT STATEMENT | | 25334 | 3958K | | 102K (1) | 00:00:05 |
| * 1 | FILTER | | | | | | |
| 2 | HASH GROUP BY | | 25334 | 3958K | 82M | 102K (1) | 00:00:05 |
| * 3 | HASH JOIN | | 506K | 77M | 81M | 95864 (1) | 00:00:04 |
| 4 | TABLE ACCESS FULL | VICTIMS | 4082K | 35M | | 5259 (1) | 00:00:01 |
| * 5 | HASH JOIN | | 904K | 130M | 40M | 79565 (1) | 00:00:04 |
| * 6 | HASH JOIN | | 456K | 34M | | 19082 (1) | 00:00:01 |
| * 7 | TABLE ACCESS FULL | TYPE_OF_COLLISION | 1 | 13 | | 3 (0) | 00:00:01 |
| 8 | TABLE ACCESS FULL | COLLISIONS | 3678K | 235M | | 19069 (1) | 00:00:01 |
| 9 | TABLE ACCESS FULL | PARTIES | 7286K | 493M | | 29872 (1) | 00:00:02 |

Predicate Information (identified by operation id):

```
" 1 - filter(MIN("V"."VICTIM_AGE")>100) "
" 3 - access("V"."PARTY_ID"="P"."ID") "
" 5 - access("P"."COLLISION_CASE_ID"="C"."CASE_ID") "
```

```
" 6 - access("C"."TYPE_OF_COLLISION_ID"="TOC"."ID")"
" 7 - filter("TOC"."DEFINITION"='Vehicle/Pedestrian')"
```

Improved plan:

| Id | Operation | Name | Rows | Bytes | TempSpc | Cost | (%CPU) | Time |
|-----|----------------------|---|-------|-------|---------|-------|--------|----------|
| 0 | SELECT STATEMENT | | 25334 | 3958K | | 86094 | (1) | 00:00:04 |
| 1 | SORT ORDER BY | | 25334 | 3958K | 82M | 86094 | (1) | 00:00:04 |
| * 2 | FILTER | | | | | | | |
| 3 | HASH GROUP BY | | 25334 | 3958K | 82M | 86094 | (1) | 00:00:04 |
| * 4 | HASH JOIN | | | | | | | |
| 5 | INDEX FAST FULL SCAN | VICTIMS_IDX_PARTY_ID_VICTIM_AGE | 506K | 77M | 81M | 72404 | (1) | 00:00:03 |
| * 6 | HASH JOIN | | 4082K | 35M | | 3063 | (1) | 00:00:01 |
| 7 | NESTED LOOPS | | 904K | 130M | 40M | 58301 | (1) | 00:00:03 |
| * 8 | INDEX RANGE SCAN | TYPE_OF_COLLISION_IDX_DEFINITION_ID | 456K | 34M | | 5002 | (1) | 00:00:01 |
| * 9 | INDEX RANGE SCAN | COLLISIONS_IDX_CASE_ID_TYPE_OF_COLLISION_ID | 1 | 13 | | 1 | (0) | 00:00:01 |
| 10 | INDEX FAST FULL SCAN | PARTIES_IDX_COLLISION_CASE_ID_ID | 456K | 29M | | 5001 | (1) | 00:00:01 |
| | | | 7286K | 493M | | 22688 | (1) | 00:00:01 |

Predicate Information (identified by operation id):

```
" 2 - filter(MIN("V"."VICTIM_AGE")>100)"
" 4 - access("V"."PARTY_ID"="P"."ID")"
" 6 - access("P"."COLLISION_CASE_ID"="C"."CASE_ID")"
" 8 - access("TOC"."DEFINITION"='Vehicle/Pedestrian')"
" 9 - access("C"."TYPE_OF_COLLISION_ID"="TOC"."ID")"
```

Query 8

Initial time: 3,5s

Optimized time: 3,24s

Explain the improvement:

We created the following 2 indexes:

```
CREATE INDEX
PARTIES_IDX_STATEWIDE_VEHICLE_TYPE_ID_VEHICLE_MAKE_VEHICLE_YEAR on
PARTIES (STATEWIDE_VEHICLE_TYPE_ID, VEHICLE_YEAR, VEHICLE_MAKE);
CREATE INDEX STATEWIDE_VEHICLE_TYPE_IDX_DEFINITION_ID on
STATEWIDE_VEHICLE_TYPE (DEFINITION, ID);
```

In this query, we access the party table over the vehicle_make, vehicle_year and the statewide_vehicle_type_id together and therefore our first index on those attributes improves the plan by replacing the TABLE ACCESS FULL with an INDEX FAST FULL SCAN. Our second index is on the statewide vehicle type table where we created an index on the id and the definition which are accessed in the same where clause. This index transforms the TABLE ACCESS FULL into an INDEX FULL SCAN. The latter means that all the needed attributes were present in the index. It only reads the index block by block and not the full table thus reducing the amount of IO.

Initial plan:

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|-------------------|------------------------|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 8935 | 305K | 30218 (2) | 00:00:02 |
| 1 | SORT ORDER BY | | 8935 | 305K | 30218 (2) | 00:00:02 |
| * 2 | FILTER | | | | | |
| 3 | HASH GROUP BY | | 8935 | 305K | 30218 (2) | 00:00:02 |
| * 4 | HASH JOIN | | 5415K | 180M | 29936 (1) | 00:00:02 |
| 5 | TABLE ACCESS FULL | STATEWIDE_VEHICLE_TYPE | 15 | 330 | 3 (0) | 00:00:01 |
| * 6 | TABLE ACCESS FULL | PARTIES | 5415K | 67M | 29919 (1) | 00:00:02 |

Predicate Information (identified by operation id):

```

2 - filter(COUNT(*)>=10)
" 4 - access("P"."STATEWIDE_VEHICLE_TYPE_ID"="SVT"."ID") "
" 6 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL AND "P"."VEHICLE_YEAR" IS NOT "
"      NULL AND "P"."VEHICLE_MAKE" IS NOT NULL) "

```

Improved plan:

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
|-----|----------------------|---|-------|-------|-------------|----------|
| 0 | SELECT STATEMENT | | 12466 | 426K | 6730 (5) | 00:00:01 |
| 1 | SORT ORDER BY | | 12466 | 426K | 6730 (5) | 00:00:01 |
| * 2 | FILTER | | | | | |
| 3 | HASH GROUP BY | | 12466 | 426K | 6730 (5) | 00:00:01 |
| * 4 | HASH JOIN | | 5415K | 180M | 6449 (1) | 00:00:01 |
| 5 | INDEX FULL SCAN | STATEWIDE_VEHICLE_TYPE_IDX_DEFINITION_ID | 15 | 330 | 1 (0) | 00:00:01 |
| * 6 | INDEX FAST FULL SCAN | PARTIES_IDX_STATEWIDE_VEHICLE_TYPE_ID_VEHICLE_MAKE_VEHICLE_YEAR | 5415K | 67M | 6434 (1) | 00:00:01 |

Predicate Information (identified by operation id):

```

2 - filter(COUNT(*)>=10)
" 4 - access("P"."STATEWIDE_VEHICLE_TYPE_ID"="SVT"."ID") "
" 6 - filter("P"."STATEWIDE_VEHICLE_TYPE_ID" IS NOT NULL AND "P"."VEHICLE_YEAR" IS NOT NULL AND "P"."VEHICLE_MAKE" IS NOT NULL) "

```

General Comments

The queries for the last milestone were more complicated and took us quite a long time, but we managed to finish them on time.

The time split between the members was almost equal and we all participated in every task during the whole semester.